

# ADK Multi-Agent System Best Practices

A deep-dive into the architectural, efficiency, and structural patterns for production-ready agents.

---

# The Core Architecture

Understanding the rationale behind multi-agent systems and the role of delegation.

# Three Pillars of ADK Modularity

## 1. Agents as Modules

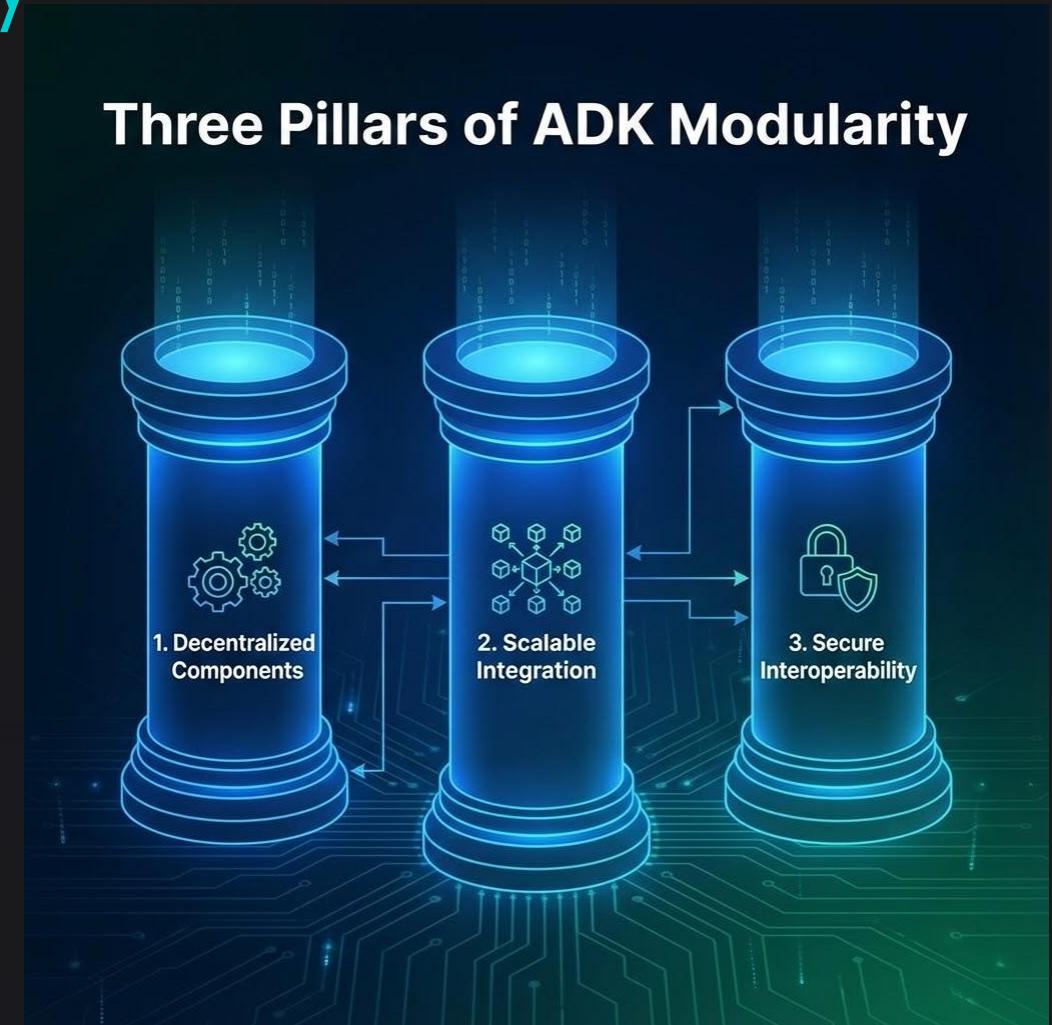
Each agent (super or sub) must be a self-contained Python package, enabling isolated development and ownership. This allows developers to work on different agents simultaneously.

## 2. Shared vs. Specific Tools

Clearly distinguish between generic tools (placed in a central `shared/` directory) and tools tightly coupled to a specific agent (placed within that agent's folder).

## 3. Explicit Interfaces

Define clear input/output contracts (data formats, expected states) between Super-Agents and their Sub-Agents to ensure reliable orchestration.



# Agent Modularity: The Foundation of Collaboration

-  **Isolate Codebase:** By making each agent a self-contained package, cross-package dependencies are minimized, reducing merge conflicts.
-  **Clear Ownership:** A single developer or team can own a single agent directory, including its code, configuration, and agent-specific tools.
-  **Easy Testing:** Allows for seamless unit testing of individual agent logic without needing to mock the entire multi-agent system.
-  **CLI Discoverability:** ADK's CLI automatically discovers agents in directories containing an `\_\_init\_\_.py` and an `agent.py` file defining a `root\_agent`.

# Rationale: Why Use Multi-Agents?

- **Complexity Management:** Decompose a large, complex problem into smaller, manageable sub-tasks. This is difficult for a single, monolithic agent.
- **Reusability & Modularity:** Specialized agents (e.g., Research, Coding) can be reused across different applications without having to rewrite their logic or tools.
- **Higher Quality Output:** A "society of mind" where experts handle specific domains leads to better reasoning, fewer hallucinations, and more precise tool use.

## LLM-Driven Delegation (Agent Transfer)

The Orchestrator's LLM dynamically routes tasks to a sub-agent using the `transfer\_to\_agent` mechanism based on the sub-agent's precise \*\*`description`\*\* field.

**Clear and distinct descriptions are vital for effective dynamic routing.**

# Best Practice: Model Selection for Scale



## Prioritize Flash

The sample uses 'gemini-2.5-flash' for all agents. This is the best practice default for multi-agent systems to minimize latency and manage operational costs.



## Cost & Latency

Since MAS involves multiple LLM calls per request, the overhead in tokens and time accumulates quickly. Flash is optimized for speed and cost-effectiveness in these scenarios.



## When to Use Pro

Reserve larger models (e.g., Gemini 2.5 Pro) for single agents responsible for complex, non-latency-sensitive tasks like final synthesis or complex multi-step reasoning.

---

# The Recommended Structure

A scalable, modular directory layout designed for the Agent Development Kit.

# Project Structure & Organization

## Clear Separation of Concerns

**Agents:** Dedicated directories for each agent's definitions and implementation logic (Orchestrator, Research, Coding).

**Shared:** Resources common to all agents, ensuring code reuse and a single source of truth for tools and utilities.

**Tests:** Separate suites for Unit tests (tool functions) and Integration tests (agent configurations).

## The Hierarchy (Simplified)

```
adk-sample/
├── agents/ (SO, RA, CA)
├── shared/
│   └── tools/
│       └── utils/
└── tests/
    ├── unit/
    └── integration/
        └── pyproject.toml
```

# Project Root Layout: A High-Level View

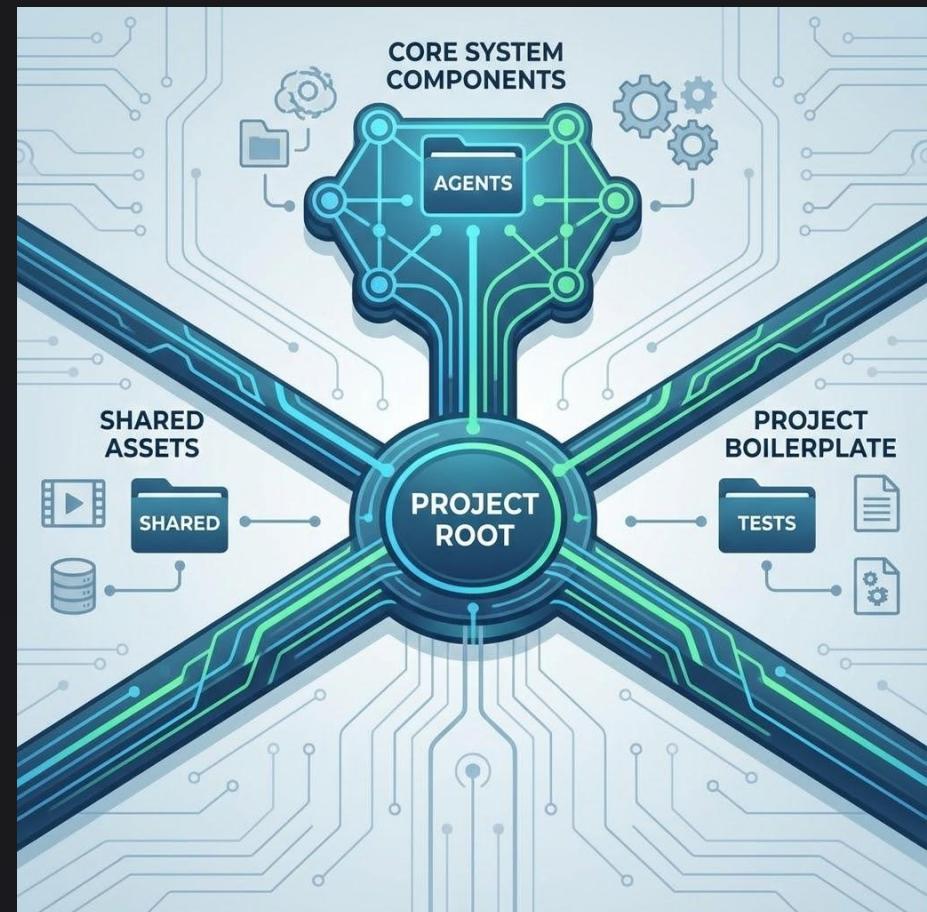
## The Top-Level Breakdown

The project root acts as the central hub, clearly dividing core system components, shared assets, and project boilerplate.

### Key Folders:

-  `agents/`: All modular agents.
-  `shared/`: Reusable tools and utilities.
-  `tests/`: Mirrored structure for unit and integration testing.

**Files:** `'.env'`, `'main.py'`, `'requirements.txt'`.



# Key Components: Agent Roles



## Super Orchestrator

Top-level agent. Delegates tasks to sub-agents. Model: 'gemini-2.5-flash'.



## Research Subagent

Specialized in info gathering. Tools:  
'brave\_search' (shared),  
'query\_knowledge\_base' (local).



## Coding Subagent

Specialized in code analysis and linting. Primary tool is 'run\_linter' (local).

# Best Practice: Unified Tool Architecture

- ⌚ **Reduce LLM Workload:** Consolidate multiple small tool functions into a single, unified tool where appropriate (e.g., a "Code Analyzer" tool that runs linter, static analysis, and formatter internally).
- ➥ **Lower Cost:** Calling one comprehensive tool instead of three separate ones significantly reduces total LLM tokens used for reasoning, leading to measurable cost and latency savings.
- ⬇️ **Structured Output:** Unified tools can provide consistent, structured outputs that the LLM finds easier to parse, improving success rate even if the output is technically longer.

# Best Practice: Context & State Management

- ⌚ **Shared Session State:** Agents exchange information and maintain conversation flow via the shared 'session.state'. This is the primary "knowledge sharing" mechanism.
- 💾 **Context Preservation:** The Orchestrator can inject relevant state (e.g., intermediate results, user tier) into the sub-agent's 'InvocationContext' during transfer to maintain the thread of thought.
- 📁 **Scoped Storage:** ADK supports state scoping ('temp:', 'session:', 'user:', 'app:') for managing data lifespan, which is critical for scalable, enterprise-ready systems.

# Tool Use: Shared vs. Local

## Shared Tools (`shared/tools/`)

- Multiple agents need the exact same functionality (e.g., search, time).
- Promotes \*\*code reusability\*\* and ensures a single source of truth.
- Best for generic utilities to simplify maintenance.

## Local Tools (`agents/agent/tools.py`)

- The tool is specific to one agent's domain (e.g., `run\_linter`).
- Tool logic is tightly coupled to the agent's role.
- Ensures \*\*modularity\*\* and keeps agent code highly self-contained.

# The Foundation: Comprehensive Testing

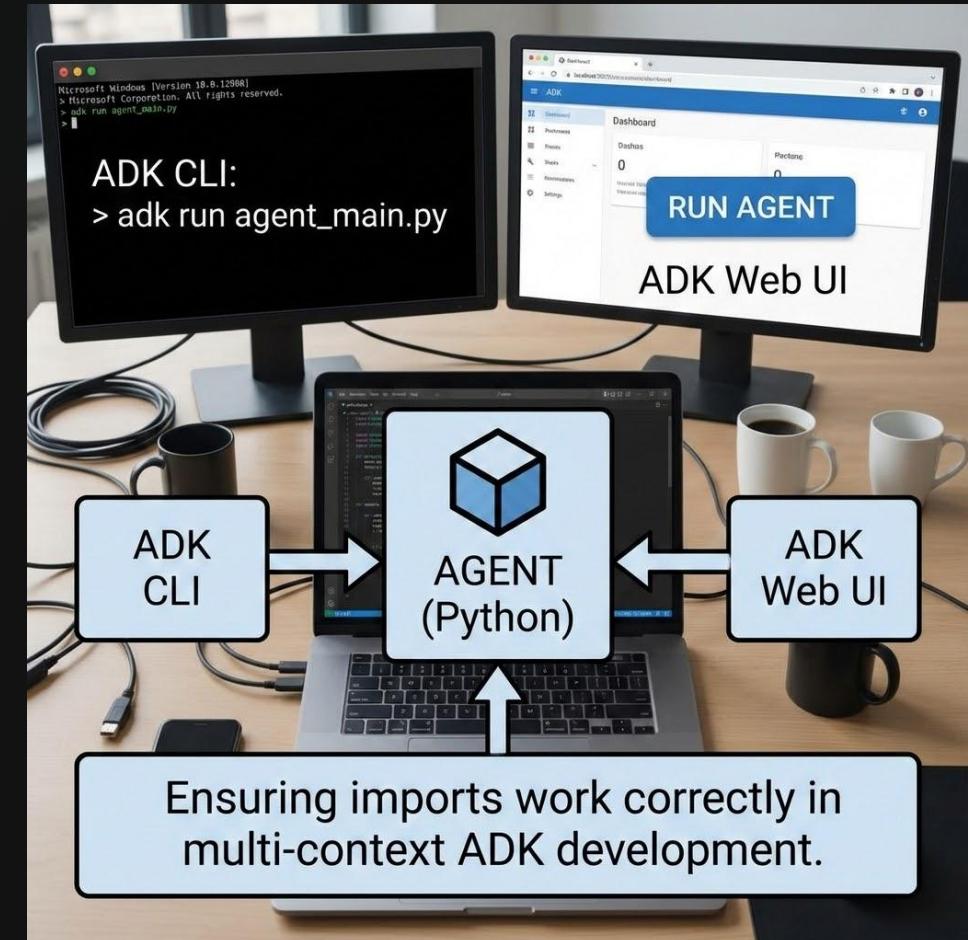
- **Unit Tests:** Verify individual tool functions in isolation, ensuring tool return values, behavior, and error handling are correct.
- ▶ **Integration Tests:** Verify agent configuration, ensuring the correct tools are attached and agent metadata (name, model) is set properly.
- ☒ **Command:** Run the entire test suite from the project root using:  
`python3 -m unittest discover tests  
-v`

# Technical Best Practice: Robust Imports

Ensuring imports work correctly is critical in multi-context ADK development. Agents must be runnable via:

- **ADK CLI:** adk run agents/X
- **ADK Web UI:** adk web agents

**The solution is to manage sys.path in agent files to ensure the project root is always resolvable.**



# Running the Agents

## Run Individual Agents (CLI)

Run the top-level orchestrator agent directly from the command line, which automatically includes its sub-agents for delegation.

```
adk run agents/super_orchestrator
```

## Run with Web UI

Launch all agents simultaneously with a user-friendly web interface for interaction and inspection (accessible at <http://127.0.0.1:8000>).

```
adk web agents
```

# Orchestration: Flat vs. Nested Hierarchy

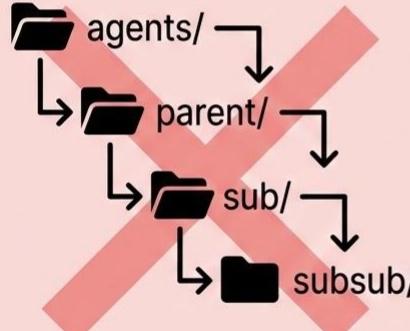
## Keep the Agent Structure Flat

**AVOID:** Deeply nested folder structures like

`agents/parent/sub/subsub/`. This leads to hard-to-navigate and tightly-coupled architectures.

**BEST PRACTICE:** Keep agents relatively flat within the `agents/` directory.

The relationship (e.g., Orchestrator calls Sub-agent) should be defined explicitly in the **code** (in the orchestrator `agent.py`) or via **configuration**, not implicitly by folder nesting.

Orchestration: Flat vs. Nested Hierarchy	
AVOID: Deeply Nested Hierarchy	BEST PRACTICE: Keep Agents Flat
 <p>Hard-to-navigate, Tightly-coupled Architecture</p>	 <p>Relatively flat within agents/  Explicit relationship defined in code/configuration</p>

# Parallel Development & Testing

 **Parallel Workstreams:** Developer A works solely in `agents/research\_subagent/` while Developer B works in `agents/coding\_subagent/`.

 **Minimal Cross-Paths:** Teams only cross paths when updating the contracts (interfaces) or modifying code within `shared/tools/`.

 **Isolated Testing:** Run unit tests for just one developer's agent without affecting the others: `pytest tests/unit/agents//`.

 **Scalability:** This modular approach ensures that adding new, specialized agents (e.g., a `deployment\_subagent/`) does not introduce complexity to existing agents.

# Questions?

Thank you for exploring the ADK Multi-Agent Best Practices.

Google ADK Documentation